

Paper 1: Li et al.

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Title Making Geo-Replicated Systems Fast as
Possible, Consistent when necessary

*10th USENIX Symposium on Operating Systems Design and
Implementation*

Authors: Cheng Li, Daniel Porto, Allen Clement, Johannes
Gehrke, Nuno Preguica, Rodrigo Rodrigues

Date: 2012

Motivation:

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 To improve user-experience, services replicate system state across geographically diverse sites.
- 2 Performance vs Consistency
 - Amazon's Dynamo - eventual consistency where state temporarily converge.
 - Yahoo PNUTS - avoids state divergence by requiring all operations that update the service state to be funneled through a primary site and thus incurring increased latency.

Overview:

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 RedBlue Consistency - Blue operations execute locally and are lazily replicated. Red operations are serialized with respect to each other and are immediately cross-site coordinated.
- 2 Conditions under which operations must be colored red or blue.
- 3 Decomposing operations into two components - a generator operation and a shadow operation.

Properties of Geo-Replicated Systems

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Low latency - Operations should proceed after contacting a small number of users.
- 2 Causality - Monotonicity of user request within session and also preserving causality across clients
- 3 State Convergence - All replicas have executed the same set of operations
- 4 All operations should return a single value.
- 5 The system should provide a set of stable histories and support for general operations.
- 6 The system should preserve a set of invariants.
- 7 Eventual Propagation

Related Work: Consistency

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Consistency level	Example systems	Immediate response	State convergence	Single value	General operations	Stable histories	Classification strategy
Strong	RSM [20, 31]	no	yes	yes	yes	yes	N/A
Timeline/snapshot	PNUTS [8], Megastore [3]	reads only	yes	yes	yes	yes	N/A
Fork	SUNDR [24]	all ops	no	yes	yes	yes	N/A
Eventual	Bayou [38], Depot [26]	all ops	yes	no	yes	yes	N/A
	Spore [12], CRDT [33]	all ops	yes	yes	no	yes	N/A
	Zeno [34], COPS [25]	weak/all ops	yes	yes	yes	no	no / N/A
	PSI [35]	cset	yes	yes	partial	yes	no
Multi	lazy repl. [19], Horus [39]	immed./causal ops	yes	yes	yes	yes	no
RedBlue	Gemini	Blue ops	yes	yes	yes	yes	yes

Table 1: Tradeoffs in geo-replicated systems and various consistency levels.

Related Work: Levels of Consistency

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Strong Consistency - Replicated systems behave like a single server that serialize all operations.
- 2 Timeline/Snapshot Consistency - There is a total order for updates to the service state but gives the option of reading a consistent but dated view of the service.
- 3 Fork Consistency - Relaxes strong consistency by allowing users to observe distinct causal histories.
- 4 Eventual Consistency - All replicas "eventually" diverge at some state.
- 5 Multi Consistency - Other systems expose multiple values from divergent branches in operation replies either directly to the client or to an application-specific conflict resolution procedure.
- 6 RedBlue Consistency - Operations have multiple consistency levels

Related Work: Other

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Consistency Rationing - Consistency guarantees associated with the data instead of the operation. Also switches consistency levels at runtime.
- 2 TACT - bounds the amount of inconsistency based on parameters like numeric errors, order errors, staleness, etc.

System Model - Assumptions

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 A distributed system with state fully replicated across k sites denoted $site_0 \dots site_{k-1}$
- 2 $S \in \mathcal{S}$ denotes a system state and $u, v \in \mathcal{O}$ a set of operations.
- 3 Initial State - S_0 . When operation u is applied it goes to state S' . So $S' = S + u$
- 4 $\forall S \in \mathcal{S}, S + u + v = S + v + u$
- 5 A state S is valid if it satisfies all these invariants.
- 6 Each u is submitted to one site which is called u 's primary site and denoted by $site(u)$.
- 7 The system later replicates u to the other sites.

RedBlue Consistency

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- RedBlue order : Given a set of operations $U = B \cup R$, where $B \cap R = \emptyset$, a RedBlue order is a partial order $O = (U, \prec)$ with the restriction that $\forall u, v \in R$ such that $u \neq v$, $u \prec v$ or $v \prec u$ (i.e. red operations are totally ordered).
- Causal Serialization : Given a site i , $O_i = (U, <)$ is an i -causal serialization (or short, a causal serialization) of RedBlue order $O = (U, \prec)$ if
 - 1 O_i is a linear extension of O (i.e. i is a total order compatible with the partial order \prec)
 - 2 for any two operations $u, v \in U$, if $site(v) = i$ and $u < v$ in O_i then $u \prec v$

RedBlue Consistency - Definition

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- RedBlue consistency : A replicated sytem is O-RedBlue consistent(or short, RedBlue consistent) if each site i applies operations according to an i -causal serialization of RedBlue order O .

Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick


Li et al.

Lloyd et al.

Bibliography

Alice in EU

Bob in US



(a) RedBlue order O of operations

Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



(b) Causal serializations of O

State Convergence

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- A RedBlue consistent system is state convergent if all causal serializations of the underlying RedBlue order O reach the same state S .

State Convergence: Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

```
1 float balance, interest = 0.05;
2 func deposit( float money ):
3     balance = balance + money;
4 func withdraw ( float money ):
5     if ( balance - money >= 0 ) then:
6         balance = balance - money;
7     else print "failure";
8 func accrueinterest():
9     float delta = balance × interest;
10    balance = balance + delta;
```

State Convergence: Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Alice in EU

Δ *deposit(20)*

Bob in US

Δ *accrueinterest()*

(a) RedBlue order O of operations issued by Alice and Bob

Alice in EU

balance:100

Δ *deposit(20)*

balance:120

Δ *accrueinterest()*

balance:126

Bob in US

balance:100

Δ *accrueinterest()*

balance:105

Δ *deposit(20)*

balance:125

\neq

State Convergence:Theorem

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Theorem

Given a RedBlue order O , if all blue operations are globally commutative then any O -RedBlue consistent system is state convergent

Replicating side effects -Generator Operation and Shadow Operation

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Generator Operation g_u - executed only at primary site against some system state S .
- 2 Shadow Operation $h_u(S)$ - executed at every site(including the primary site)

Replicating side effects - Defining shadow operations

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Correct Generator/ Shadow Operations : The decomposition of operation u into generator and shadow operations is correct if for all states S , the generator operation g_u has no effect and the generated shadow operation $h_u(S)$ has the same effect as u , i.e., for any state S : $S + g_u = S$ and $S + h_u(S) = S + u$

Shadow Banking and Invariants - Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

```
1  func deposit' ( float money ):  
2      balance = balance + money;  
3  func withdrawAck' ( float money ):  
4      balance = balance - money;  
5  func withdrawFail' ():  
6      /* no-op */  
7  func accrueinterest' ( float delta ):  
8      balance = balance + delta;
```

Shadow Banking and Invariants - Example


Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick


Li et al.

Lloyd et al.

Bibliography



(a) RedBlue order O of banking shadow operations



Shadow Banking and Invariants - Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Invariant Safe - Shadow operation $h_u(S)$ is invariant safe if for all valid states S and S' , the state $S' + h_u(S)$ is also valid.

Theorem

If all shadow operations are correct and all blue shadow operations are invariant safe and globally commutative, then for any execution of that system that is RedBlue consistent, no site is ever in an invalid state.

What can be blue? What can be red?

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

The procedure for deciding which shadow operations can be blue or must be red if a RedBlue consistent system is to provide both state convergence and invariant preservation:

- 1 For any pair of non-commutative shadow operations u and v , label both u and v red.
- 2 For any shadow operation u that may result in an invariant being violated, label u red.
- 3 Label all non-red shadow operations blue.

Shadow Banking and Invariants - Example

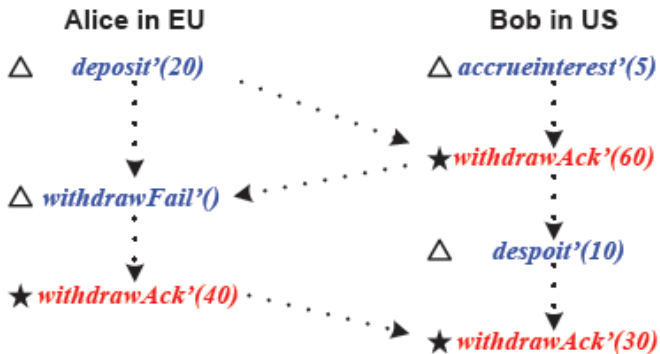
Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



(a) RedBlue order O of banking shadow operations

Shadow Banking and Invariants - Example


Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



(b) Convergent and invariant preserving causal serializations of O

Case Studies

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- TPC-W shopping cart benchmark
- RUBiS auction benchmark
- Quoddy social networking application

Two main tasks :

- Decomposing the application into a generator and shadow operation
- Labeling the shadow operations appropriately

Original application to RedBlue Consistent

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Application	Original				RedBlue consistent extension					
	user requests	transactions			LOC	shadow operations				LOC changed
		total	read-only	update		blue no-op	blue update	red	LOC	
TPC-W	14	20	13	7	9k	13	14	2	2.8k	429
RUBiS	26	16	11	5	9.4k	11	7	2	1k	180
Quoddy	13	15	11	4	15.5k	11	4	0	495	251

Table 2: Original applications and the changes needed to make them RedBlue consistent.

TPC-W

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Serves 14 different user requests such as browsing, searching, adding products to a shopping cart or placing an order.
- Each user request generates one to four transactions that access state stored across eight different tables.
- Shopping cart can be shared by multiple users across multiple sessions.

TPC-W - Writing TPC-W generator and shadow operations

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

```
1  doBuyConfirm(canId) {  
2    beginTxn();  
3    cart = exec(SELECT * FROM cartTb WHERE cId=canId);  
4    cost = computeCost(cart);  
5    orderId = getUniqueId();  
6    exec(INSERT INTO orderTb VALUES(orderId, cart.item.id, cart.item.qty  
      , cost));  
7    item = exec(SELECT * FROM itemTb WHERE id=can.item.id);  
8    if item.stock - cart.item.qty < 10 then:  
9      delta = item.stock - cart.item.qty + 21;  
10     if delta > 0 then:  
11       exec(UPDATE itemTb SET item.stock += delta);  
12     else rollback();  
13   else exec(UPDATE itemTb SET item.stock -= cart.item.qty);  
14   exec(DELETE FROM cartContentTb WHERE cId=canId AND id=  
      cart.item.id);  
15   commit(); }
```

(a) Original transaction that commits changes to database.

TPC-W - Writing TPC-W generator and shadow operations

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

```
1  doBuyConfirmGenerator(carId) {
2    sp = getScratchpad();
3    sp.beginTxn();
4    car = sp.exec(SELECT * FROM cartTb WHERE cId=carId);
5    cost = computeCost(car);
6    orderId = getUniqueId();
7    sp.exec(INSERT INTO orderTb VALUES (orderId, car.item.id,
8      car.item.qty, cost));
9    item = sp.exec(SELECT * FROM itemTb WHERE id=car.item.id);
10   if item.stock - car.item.qty < 10 then:
11     delta = item.stock - car.item.qty + 21;
12     if delta > 0 sp.exec(UPDATE itemTb SET item.stock += delta);
13     else sp.discard(); return;
14   else sp.exec(UPDATE itemTb SET item.stock -= car.item.qty);
15   sp.exec(DELETE FROM cartTb WHERE cId=carId AND id=car.item.id);
16   LTS = getCommitOrder();
17   sp.discard();
18   if replenished return (doBuyConfirmIncr' (orderId, carId,
19     car.item.id, car.item.qty, cost, delta, LTS));
20   else return (doBuyConfirmDecr' (orderId, carId, car.item.id,
21     car.item.qty, cost, LTS)); }
```

(b) Generator operation that manipulates data via a private *scratchpad*.

TPC-W - Writing TPC-W generator and shadow operations

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

```
1 doBuyConfirmIncr' (orderid, cartid, uid, qty, cost, delta, LJS) {  
2   exec(INSERT INTO orderTb VALUES (orderid, uid, qty, cost, LJS));  
3   exec(UPDATE itemTb SET item.stock += delta);  
4   exec(UPDATE itemTb SET item.ljs = LJS WHERE item.ljs < LJS);  
5   exec(UPDATE cartContentTb SET flag = TRUE WHERE id = uid AND  
      cid = cartid AND ljs <= LJS);
```

(c) Shadow doBuyConfirmIncr' (Blue) that replenishes the stock value.

```
1 doBuyConfirmDecr' (orderid, cartid, uid, qty, cost, LJS) {  
2   exec(INSERT INTO orderTb VALUES (orderid, uid, qty, cost, LJS));  
3   exec(UPDATE itemTb SET item.stock -= qty);  
4   exec(UPDATE itemTb SET item.ljs = LJS WHERE item.ljs < LJS);  
5   exec(UPDATE cartContentTb SET flag = TRUE WHERE id = uid AND  
      cid = cartid AND ljs <= LJS);
```

(d) Shadow doBuyConfirmDecr' (Red) that decrements the stock value.

Evaluation-Experimental Setup

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Experiments were run on Amazon EC2 using 5 virtual machine instances located in 5 sites - US east(UE), US west(UW), Ireland(IE), Brazil(BR) and Singapore(SG).
- Each VM has 8 virtual cores and 15 GB of RAM. VMs run Debian 6(Squeeze) 64 bit, MYSQL 5.5.18, Tomcat 6.0.35 and Sun Java SDK 1.6.

Evaluation-TPC-W

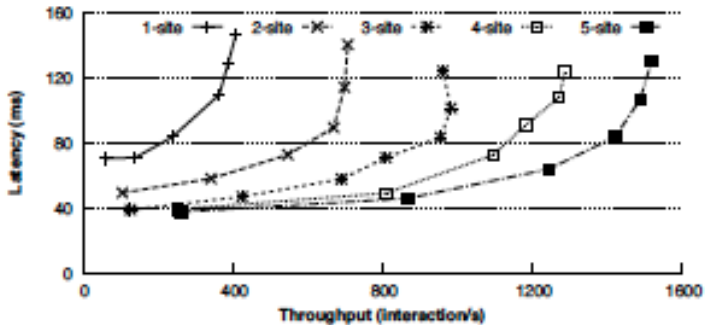
Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



(a) TPC-W shopping mix

Paper 2: Lloyd et al.

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Stronger Semantics for Low-Latency Geo-Replicated Storage

Proceedings of the 10th USENIX Symposium on Networked Systems Design and Implementation (NSDI13)

Wyatt Lloyd, Michael J. Freedman, Michael Kaminsky, and
David G. Andersen

April 2013

Motivation

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Consider an example from Facebook

Motivation

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Consider an example from Facebook
- Joe performs 2 actions:
 - Defriend boss.
 - Make snarky comment about boss and promise to quit.
- In an **eventually** consistent system, these can be viewed in the wrong order for a period time, earning Joe his first unemployment check.
- This won't happen with **causal** consistency.

Consistency - Causal versus Eventual

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Eventual Consistency
 - At some point, the replicas will converge
 - Until then, no guarantees
 - Joe's boss might see his whining.

Consistency - Causal versus Eventual

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

■ Eventual Consistency

- At some point, the replicas will converge
- Until then, no guarantees
- Joe's boss might see his whining.

■ Causal Consistency

- At a single processor, serial order of events determines causal order
- Reads are causally ordered after their writes across processors
- Transitive closure of these two properties
- Joe is safe to gripe.

Overview

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- We modify an eventually consistent system to get a causally consistent one
- Extend Cassandra to get Eiger
- Take slight hit in **throughput** to get stronger version of **consistency**
- Maintain low latency

Contributions of the paper

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Eiger
 - Low Latency
 - High throughput (slightly lower than Cassandra)
 - Causal Consistency (rather than eventual as in Cassandra)
- Read Only Algorithm
- Write Only Algorithm

Background

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

■ Cassandra

- Designed by Facebook to search inbox messages (has since been replaced by Hbase for better scalability)
- Other companies are still using Cassandra (e.g., Netflix and Apple)
- Allows Eventual or Strong Consistency
- If you want low latency, Eventual is the only option

Eiger: Column Family Data Model

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Hierarchical Column families: 3 tiers
- Example:
 - Super Family = Association
 - Family = Friends
 - Column = Alice

User Data			Associations				
			Friends			Likes	
	ID	Town	Alice	Bob	Carol	NSDI	SOSP
Alice	1337	NYC	-	3/2/11	9/2/12	9/1/12	-
Bob	2664	LA	3/2/11	-	-	-	-
⋮							

Figure 1: An example use of the column-family data model for a social network setting.

Eiger: Operations

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Cassandra has only 3 operations (key determines row)
 - Insert(table, key, rowMutation)
 - Get(table, key, columnName)
 - Delete(table, key, columnName)

Eiger: Operations

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Cassandra has only 3 operations (key determines row)
 - Insert(table, key, rowMutation)
 - Get(table, key, columnName)
 - Delete(table, key, columnName)
- Eiger adds more complex transactions
 - Batch Mutate: insert or delete multiple keys as a group of independent operations
 - Atomic Mutate: insert or delete multiple keys in a single atomic batch
 - Multiget Slice: read multiple columns/keys
 - Multiget Slice by Time

Eiger: Client Library

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Manages read and write algorithms
- Tracks causality by maintaining 1-hop dependencies
- For an atomic write, no write is applied in a cluster until after all the dependencies have been applied. **This is necessary for causal consistency.**
- Dependencies are on operations rather than values (as in COPS).

Eiger: One-Hop Dependencies

Consistency in
the Cloud II


Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Example: For w_8 , only w_6 and w_5 are stored.



Eiger: Read Algorithm

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Non-blocking: other processes can read
- Partition Tolerant
- 1 or 2 rounds of non-blocking parallel reads

Eiger: Read Algorithm


Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick


Li et al.

Lloyd et al.

Bibliography



(a) One Round Sufficient



(b) Two Rounds Needed

Eiger: Read Algorithm

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Step 1: Find latest lower bound, L
- Step 2: Find lowest upper bound that is higher than L
If all upper bounds are higher than L then we are done. If not, continue.
- Step 3: For any reads that had an upper bound lower than L , read again.

Eiger: Write Algorithm

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Atomically write set of keys
- Lock free (and thus low latency)
- Does not block concurrent reads

Eiger: Write Algorithm

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Three parties involved

- Client - requests write to many servers
- Coordinator server - a randomly chosen server from list of transaction destinations
- Cohort server - all other destinations

Eiger: Write Algorithm

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Three parties involved

- Step 1: Break up transaction into a separate request for each destination.
- Step 2: Randomly choose coordinator and send transactions to all destinations.
- Step 3: Cohort servers send notification to coordinator.
- Step 4: Coordinator runs dependency checks
- Step 5: 2 Phase commit (to ensure all-or-nothing)
- Step 6: Coordinator sends ACK to client

Eiger: Write Algorithm


Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



Example:

Evaluation

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Comparison to Cassandra:

- Within 7% of throughput Using Facebook data

	Ops/sec	Keys/sec	Columns/sec
Cassandra	23,657	94,502	498,239
Eiger	22,088	88,238	466,844
Eiger All Txns	22,891	91,439	480,904
Max Overhead	6.6%	6.6%	6.3%

Evaluation

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Comparison to Cassandra:

- Cassandra had median latency of 0.38ms, whereas Eiger had 0.78ms
- Thus, there is about a factor of 2 difference
- Strongly Consistent Cassandra had 85.21ms latency

Evaluation

Consistency in
the Cloud II


Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Almost linear scaling:



Conclusion: Relationship between 2 Papers

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

■ Similarities

- Geo-replicated Systems
- Provide improvement in consistency

■ Differences

- Li: Latency versus Consistency.
 - Separate operations into those that must be consistent and those that need not.
- Lloyd: Consistency versus Throughput (requiring low latency)
 - Extend previous system to offer causal consistency with small hit to throughput.

The End

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Thank you!

Bibliography

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- **Stronger Semantics for Low-Latency Geo-Replicated Storage**, *Proceedings of the 10th USENIX Symposium on Networked Systems Design and Implementation (NSDI13)*, Wyatt Lloyd, Michael J. Freedman, Michael Kaminsky, and David G. Andersen, April 2013
- **Out in the Open: The Abandoned Facebook Tech That Now Helps Power Apple**, *www.wired.com* , Klint Finley, Aug. 4, 2014
- **A Short Primer on Causal Consistency**, *;login: The Usenix Magazine Volume 38, Number 4*, Wyatt Lloyd, Michael J. Freedman, Michael Kaminsky, and David G. Andersen, , August 2013
- **Netflix's Viewing Data: How We Know Where You Are in House of Cards**, <http://techblog.netflix.com/search/label/Cassandra> January 27, 2015