

Paper 1: Li et al.

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Title Making Geo-Replicated Systems Fast as
Possible, Consistent when necessary
*10th USENIX Symposium on Operating Systems Design and
Implementation*

Authors

Date

Motivation:

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 To improve user-experience, services replicate system state across geographical diverse sites.
- 2 Performance vs Consistency
 - Amazon's Dynamo - eventual consistency where state temporarily converge.
 - Yahoo PNUTS - avoids state divergence by requiring all operations that update the service state to be funneled through a primary site and thus incurring increased latency.

Overview:

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 RedBlue Consistency - Blue operations execute locally and are lazily replicated. Red operations are serialized with respect to each other and are immediately cross-site coordinated.
- 2 Conditions under which operations must be colored red or blue.
- 3 Decomposing operations into two components - a generator operation and a shadow operation.

Properties of Geo-Replicated Systems

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Low latency - Operations should proceed after contacting a small number of users.
- 2 Causality - Monotonicity of user request within session and also preserving causality across clients
- 3 State Convergence - All replicas have executed the same set of operations
- 4 All operations should return a single value.
- 5 The system should provide a set of stable histories and support for general operations.
- 6 The system should preserve a set of invariants.
- 7 Eventual Propagation

Related Work: Consistency

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Consistency level	Example systems	Immediate response	State convergence	Single value	General operations	Stable histories	Classification strategy
Strong	RSM [20, 31]	no	yes	yes	yes	yes	N/A
Timeline/snapshot	PNUTS [8], Megastore [3]	reads only	yes	yes	yes	yes	N/A
Fork	SUNDR [24]	all ops	no	yes	yes	yes	N/A
Eventual	Bayou [38], Depot [26]	all ops	yes	no	yes	yes	N/A
	Spore [12], CRDT [33]	all ops	yes	yes	no	yes	N/A
	Zeno [34], COPS [25]	weak/all ops	yes	yes	yes	no	no / N/A
	PSI [35]	cset	yes	yes	partial	yes	no
Multi	lazy repl. [19], Horus [39]	immed./causal ops	yes	yes	yes	yes	no
RedBlue	Gemini	Blue ops	yes	yes	yes	yes	yes

Table 1: Tradeoffs in geo-replicated systems and various consistency levels.

Related Work: Levels of Consistency

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Strong Consistency - Replicated systems behave like a single server that serialize all operations.
- 2 Timeline/Snapshot Consistency - There is a total order for updates to the service state but gives the option of reading a consistent but dated view of the service.
- 3 Fork Consistency - Relaxes strong consistency by allowing users to observe distinct casual histories.
- 4 Eventual Consistency - All replicas "eventually" diverge at some state.
- 5 Multi Consistency - Other systems expose multiple values from divergent branches in operation replies either directly to the client or to an application-specific conflict resolution procedure.
- 6 RedBlue Consistency - Operations have multiple consistency levels

Related Work: Other

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Consistency Rationing - Consistency guarantees associated with the data instead of the operation. Also switches consistency levels at runtime.
- 2 TACT - bounds the amount of inconsistency based on parameters like numeric errors, order errors, staleness etc.

System Model - Assumptions

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 A distributed system with state fully replicated across k sites denoted $site_0 \dots site_{k-1}$
- 2 $s \in S$ denotes a system state and $u, v \in O$ a set of operations.
- 3 Initial State - S_0 . When operation u is applied it goes to state S' . So $S' = S + u$
- 4 $\forall S \in S, S + u + v = S + v + u$
- 5 A state S is valid if it satisfies all these invariants.
- 6 Each u is submitted to one site which is called u 's primary site and denoted by $site(u)$.
- 7 The system later replicates u to the other sites.

RedBlue Consistency

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- RedBlue order : Given a set of operations $U = B \cup R$, where $B \cap R = \emptyset$, a RedBlue order is a partial order $O = (U, \prec)$ with the restriction that $\forall u, v \in R$ such that $u \neq v$, $u \prec v$ or $v \prec u$ (i.e. red operations are totally ordered).
- Causal Serialization : Given a site i , $O_i = (U, <)$ is an i -causal serialization (or short, a causal serialization) of RedBlue order $O = (U, \prec)$ if
 - 1 O_i is a linear extension of O (i.e, $<$ is a total order compatible with the partial order \prec)
 - 2 for any two operations $u, v \in U$, if $site(v) = i$ and $u < v$ in O_i then $u \prec v$

RedBlue Consistency - Definition

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- RedBlue consistency : A replicated system is O-RedBlue consistent (or short, RedBlue consistent) if each site i applies operations according to an i -causal serialization of RedBlue order O .

Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick


Li et al.

Lloyd et al.

Bibliography

Alice in EU

Bob in US



(a) RedBlue order O of operations

Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



(b) Causal serializations of O

State Convergence

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- A RedBlue consistent system is state convergent if all causal serializations of the underlying RedBlue order O reach the same state S .

State Convergence: Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

```
1 float balance, interest = 0.05;
2 func deposit( float money ):
3     balance = balance + money;
4 func withdraw ( float money ):
5     if ( balance - money >= 0 ) then:
6         balance = balance - money;
7     else print "failure";
8 func accrueinterest():
9     float delta = balance × interest;
10    balance = balance + delta;
```

State Convergence: Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Alice in EU

Δ *deposit(20)*

Bob in US

Δ *accrueinterest()*

(a) RedBlue order O of operations issued by Alice and Bob

Alice in EU

balance:100

Δ *deposit(20)*

balance:120

Δ *accrueinterest()*

balance:126

Bob in US

balance:100

Δ *accrueinterest()*

balance:105

Δ *deposit(20)*

balance:125

\neq

State Convergence:Theorem

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Theorem

Given a RedBlue order O , if all blue operations are globally commutative then any O -RedBlue consistent system is state convergent

Replicating side effects -Generator Operation and Shadow Operation

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Generator Operation g_u - executed only at primary site against some system state S .
- 2 Shadow Operation $h_u(S)$ - executed at every site(including the primary site)

Replicating side effects - Defining shadow operations

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Correct Generator/ Shadow Operations : The decomposition of operation u into generator and shadow operations is correct if for all states S , the generator operation g_u has no effect and the generated shadow operation $h_u(S)$ has the same effect as u , i.e., for any state S : $S + g_u = S$ and $S + h_u(S) = S + u$

Replicating side effects -Revisiting RedBlue consistency

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Given a site i , $O_i = (U \cup V_i, <)$ is an i -causal serialization of RedBlue order $O = (U, \prec)$ if
 - O_i is a total order;
 - $(U, <)$ is a linear extension of O ;
 - For any $h_v(S) \in U$ generated by $g_v \in V_i$, S is the state obtained after applying the sequence of shadow operations preceeding g_v in O_i ;
 - For any $g_v \in V_i$ and $h_u(S) \in U$, $h_u(S) < g_v$ in O_i iff $h_u(S) \prec h_v(S')$ in O .

Shadow Banking and Invariants - Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

```
1  func deposit' ( float money ):  
2      balance = balance + money;  
3  func withdrawAck' ( float money ):  
4      balance = balance - money;  
5  func withdrawFail' ():  
6      /* no-op */  
7  func accrueinterest' ( float delta ):  
8      balance = balance + delta;
```

Shadow Banking and Invariants - Example


Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick


Li et al.

Lloyd et al.

Bibliography



(a) RedBlue order O of banking shadow operations



Shadow Banking and Invariants - Example

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Invariant Safe - Shadow operation $h_u(S)$ is invariant safe if for all valid states S and S' , the state $S' + h_u(S)$ is also valid.

Theorem

If all shadow operations are correct and all blue shadow operations are invariant safe and globally commutative, then for any execution of that system that is RedBlue consistent, no site is ever in an invalid state.

What can be blue? What can be red?

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

The procedure for deciding which shadow operations can be blue or must be red if a RedBlue consistent system is to provide both state convergence and invariant preservation:

- 1 For any pair of non-commutative shadow operations u and v , label both u and v red.
- 2 For any shadow operation u that may result in an invariant being violated, label u red.
- 3 Label all non-red shadow operations blue.

Shadow Banking and Invariants - Example


Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



(a) RedBlue order O of banking shadow operations

Shadow Banking and Invariants - Example

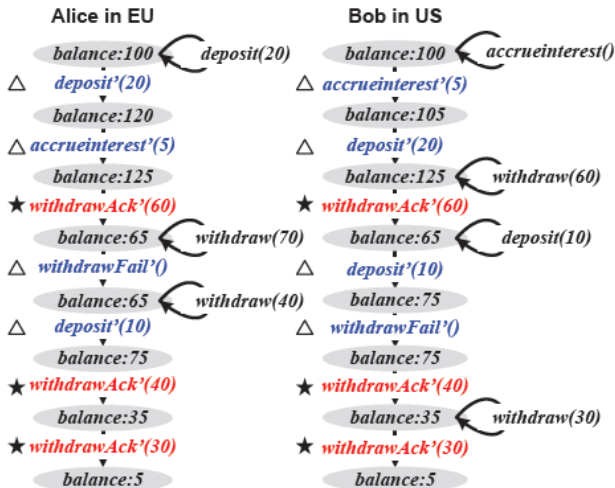
Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



(b) Convergent and invariant preserving causal serializations of O

Gemini Design and Implementation - Prototype

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 It consisted of 10K lines of java code and uses MySQL as its storage back-end
- 2 Each Gemini site consists of :
 - a storage engine
 - a proxy server
 - a concurrency coordinator
 - a data writer
- 3 The single site is replicated across multiple sites.

Gemini Design and Implementation - Basic Flow

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 User issues request to a proxy server located at the closest site.
- 2 The proxy server processes the request by executing an appropriate application transaction which is implemented as a single Gemini operation.
- 3 Storage Engine - Relational Database
- 4 Scratchpad Operations - Temporary tables

Gemini Design and Implementation - Basic Flow

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 User issues request to a proxy server located at the closest site.
- 2 The proxy server processes the request by executing an appropriate application transaction which is implemented as a single Gemini operation.
- 3 Storage Engine - Relational Database
- 4 Scratchpad Operations - Temporary tables

Gemini Design and Implementation - Failure Handling

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- 1 Isolated Component Failure - Standard state machine replication techniques can be employed to make each component robust.
- 2 Site Failure - A fault tolerance consensus protocol like Paxos can be used.
- 3 Operation Propagation - This can be addressed by using standard techniques for exchanging causal logs or reliable multicast.
- 4 Cross-session monotonicity - This can be addressed by allowing the user to specify a "last read" version when starting a new session or requiring the user to cache all relevant requests in order to reply to them when connecting to a new site.

Case Studies

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- TPC-W shopping cart benchmark
- RUBiS auction benchmark
- Quoddy social networking application

Two main tasks :

- Decomposing the application into a generator and shadow operation
- Labeling the shadow operations appropriately

Original application to RedBlue Consistent

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Application	Original					RedBlue consistent extension				
	user requests	transactions			LOC	shadow operations				LOC changed
		total	read-only	update		blue no-op	blue update	red	LOC	
TPC-W	14	20	13	7	9k	13	14	2	2.8k	429
RUBiS	26	16	11	5	9.4k	11	7	2	1k	180
Quoddy	13	15	11	4	15.5k	11	4	0	495	251

Table 2: Original applications and the changes needed to make them RedBlue consistent.

TPC-W

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Serves 14 different user requests such as browsing, searching, adding products to a shopping cart or placing an order.
- Each user request generates one to four transactions that access state stored across eight different tables.
- Shopping cart can be shared by multiple users across multiple sessions.

TPC-W - Writing TPC-W generator and shadow operations

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

```
1  doBuyConfirm(canId) {  
2    beginTxn();  
3    cart = exec(SELECT * FROM cartTb WHERE cId=canId);  
4    cost = computeCost(cart);  
5    orderId = getUniqueId();  
6    exec(INSERT INTO orderTb VALUES(orderId, cart.item.id, cart.item.qty  
      , cost));  
7    item = exec(SELECT * FROM itemTb WHERE id=can.item.id);  
8    if item.stock - cart.item.qty < 10 then:  
9      delta = item.stock - cart.item.qty + 21;  
10     if delta > 0 then:  
11       exec(UPDATE itemTb SET item.stock += delta);  
12     else rollback();  
13   else exec(UPDATE itemTb SET item.stock -= cart.item.qty);  
14   exec(DELETE FROM cartContentTb WHERE cId=canId AND id=  
      cart.item.id);  
15   commit(); }
```

(a) Original transaction that commits changes to database.

TPC-W - Writing TPC-W generator and shadow operations

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

```
1  doBuyConfirmGenerator(carId) {
2    sp = getScratchpad();
3    sp.beginTxn();
4    car = sp.exec(SELECT * FROM cartTb WHERE cId=carId);
5    cost = computeCost(car);
6    orderId = getUniqueId();
7    sp.exec(INSERT INTO orderTb VALUES (orderId, car.item.id,
8      car.item.qty, cost));
9    item = sp.exec(SELECT * FROM itemTb WHERE id=car.item.id);
10   if item.stock - car.item.qty < 10 then:
11     delta = item.stock - car.item.qty + 21;
12     if delta > 0 sp.exec(UPDATE itemTb SET item.stock += delta);
13     else sp.discard(); return;
14   else sp.exec(UPDATE itemTb SET item.stock -= car.item.qty);
15   sp.exec(DELETE FROM cartTb WHERE cId=carId AND id=car.item.id);
16   LTS = getCommitOrder();
17   sp.discard();
18   if replenished return (doBuyConfirmIncr' (orderId, carId,
19     car.item.id, car.item.qty, cost, delta, LTS));
20   else return (doBuyConfirmDecr' (orderId, carId, car.item.id,
21     car.item.qty, cost, LTS)); }
```

(b) Generator operation that manipulates data via a private *scratchpad*.

TPC-W - Writing TPC-W generator and shadow operations

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

```
1 doBuyConfirmIncr' (orderid, cartid, uid, qty, cost, delta, LJS) {  
2   exec(INSERT INTO orderTb VALUES (orderid, uid, qty, cost, LJS));  
3   exec(UPDATE itemTb SET item.stock += delta);  
4   exec(UPDATE itemTb SET item.ljs = LJS WHERE item.ljs < LJS);  
5   exec(UPDATE cartContentTb SET flag = TRUE WHERE id = uid AND  
      cid = cartid AND ljs <= LJS);
```

(c) Shadow doBuyConfirmIncr' (Blue) that replenishes the stock value.

```
1 doBuyConfirmDecr' (orderid, cartid, uid, qty, cost, LJS) {  
2   exec(INSERT INTO orderTb VALUES (orderid, uid, qty, cost, LJS));  
3   exec(UPDATE itemTb SET item.stock -= qty);  
4   exec(UPDATE itemTb SET item.ljs = LJS WHERE item.ljs < LJS);  
5   exec(UPDATE cartContentTb SET flag = TRUE WHERE id = uid AND  
      cid = cartid AND ljs <= LJS);
```

(d) Shadow doBuyConfirmDecr' (Red) that decrements the stock value.

Evaluation-Experimental Setup

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Experiments were run on Amazon EC2 using 5 virtual machine instances located in 5 sites - US east(UE), US west(UW), Ireland(IE), Brazil(BR) and Singapore(SG).
- Each VM has 8 virtual cores and 15 GB of RAM. VMs run Debian 6(Squeeze) 64 bit, MYSQL 5.5.18, Tomcat 6.0.35 and Sun Java SDK 1.6.

Evaluation-Experimental Setup

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

	UE	UW	IE	BR	SG
UE	0.4 ms 994 Mbps	85 ms 164 Mbps	92 ms 242 Mbps	150 ms 53 Mbps	252 ms 86 Mbps
UW		0.3 ms 975 Mbps	155 ms 84 Mbps	207 ms 35 Mbps	181 ms 126 Mbps
IE			0.4 ms 996 Mbps	235 ms 54 Mbps	350 ms 52 Mbps
BR				0.3 ms 993 Mbps	380 ms 65 Mbps
SG					0.3 ms 993 Mbps

Table 3: Average round trip latency and bandwidth between Amazon sites.

Evaluation-Microbenchmark

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Each user issues requests accessing a random record from a MySQL database.
- Each request maps to a single shadow operation
- A request is blue if it maps to a blue shadow operation and red otherwise
- Dataset consists of 10 tables, each initialized with 1000000 records, each record has 1 text and 4 integer attributes. The total size of the dataset is 1.0 GB.

Evaluation-Microbenchmark


Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick


Li et al.

Lloyd et al.



Bibliography



(a) Blue request latency for all users as number of sites increases



(b) Red request latency for all users as number of sites increases



Evaluation-Peak Throughput


Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



Evaluation-TPC-W



Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



Conclusion

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Evaluation-TPC-W


Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography



	TPC-W shopping	
	Original	Gemini
Thput. (inter/s)	409	386
Avg. latency	14 ms	15 ms

Paper 2: Lloyd et al.

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Stronger Semantics for Low-Latency Geo-Replicated Storage

Proceedings of the 10th USENIX Symposium on Networked Systems Design and Implementation (NSDI13)

Wyatt Lloyd, Michael J. Freedman, Michael Kaminsky, and
David G. Andersen

April 2013

Main Idea

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Main Idea

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Take slight hit in throughput to get stronger version of consistency

Main Idea

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Take slight hit in throughput to get stronger version of consistency
- Causal Consistency Instead of Eventual Consistency (causal is stronger)

Main Idea

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Take slight hit in throughput to get stronger version of consistency
- Causal Consistency Instead of Eventual Consistency (causal is stronger)
- We require low latency

Main Idea

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Take slight hit in throughput to get stronger version of consistency
- Causal Consistency Instead of Eventual Consistency (causal is stronger)
- We require low latency
- Extend previous systems: Cassandra and COPS

Contributions

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Contributions

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Eiger
 - Low Latency
 - High throughput (slightly lower than Cassandra)
 - Causal Consistency (rather than eventual as in Cassandra)

Contributions

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Eiger
 - Low Latency
 - High throughput (slightly lower than Cassandra)
 - Causal Consistency (rather than eventual as in Cassandra)
- Read Only Algorithm

Contributions

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Eiger
 - Low Latency
 - High throughput (slightly lower than Cassandra)
 - Causal Consistency (rather than eventual as in Cassandra)
- Read Only Algorithm
- Write Only Algorithm

Background

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Background

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Cassandra
 - Eventual Consistency

Background

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Cassandra
 - Eventual Consistency
- COPS

Consistency - Causal versus Eventual

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Consistency - Causal versus Eventual

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

■ p1

Column Family Data Model

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Column Family Data Model

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

■ p1

Column Family Data Model

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- p1
- p2

Column Family Data Model

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- p1
- p2
- p3

Eiger

Consistency in the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Eiger

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

■ p1

Eiger

Consistency in the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

■ p1

■ p2

Eiger

Consistency in the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- p1
- p2
- p3

Evaluation

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Evaluation

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Versus Cassandra
 - Within 7% of throughput Using Facebook-like data
 - Ops/sec
 - Keys/sec
 - Columns/sec

Evaluation

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- Versus Cassandra
 - Within 7% of throughput Using Facebook-like data
 - Ops/sec
 - Keys/sec
 - Columns/sec
- Versus COPS

Follow Up Research

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Follow Up Research

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

■ p1

Ideas for Future Research

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

Ideas for Future Research

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

■ p1

Bibliography

Consistency in
the Cloud II

Satabdi
Aditya and
Shannon
Harwick

Li et al.

Lloyd et al.

Bibliography

- **Stronger Semantics for Low-Latency Geo-Replicated Storage**, *Proceedings of the 10th USENIX Symposium on Networked Systems Design and Implementation (NSDI13)*, Wyatt Lloyd, Michael J. Freedman, Michael Kaminsky, and David G. Andersen, April 2013